

PC Hardware and x86

Dr. Jun Zheng
CSE325 Principles of Operating
Systems
8/23/2019



A PC



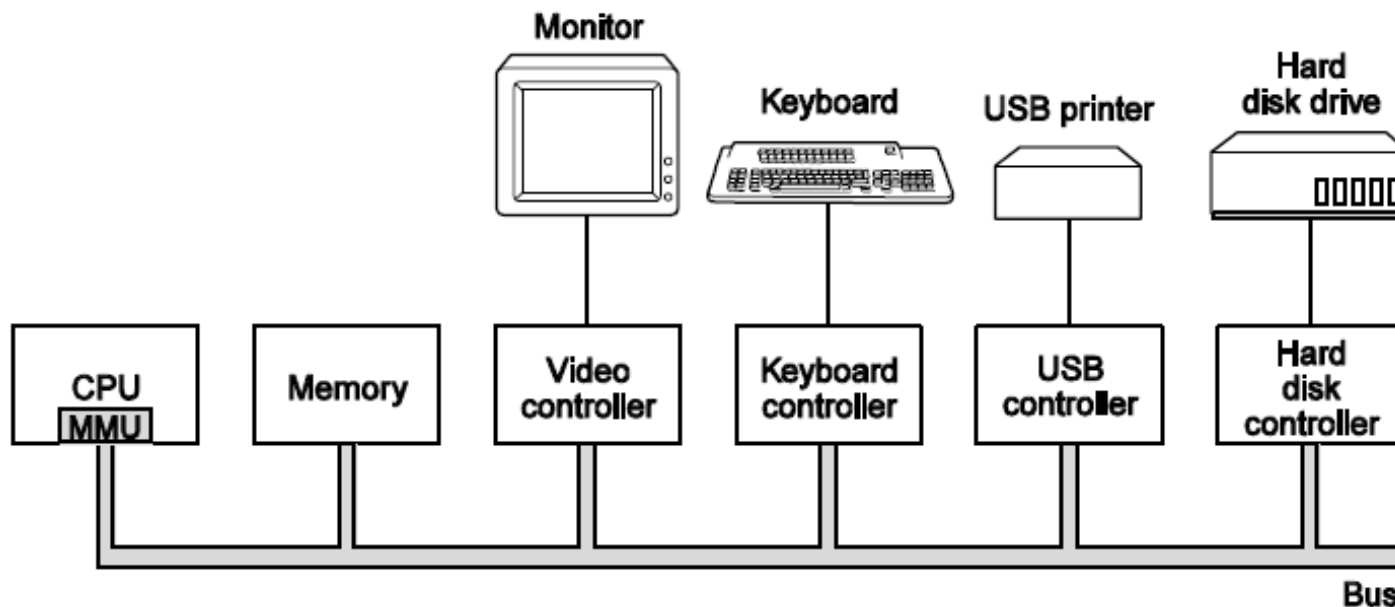
□ How to make it do something useful?

Outline

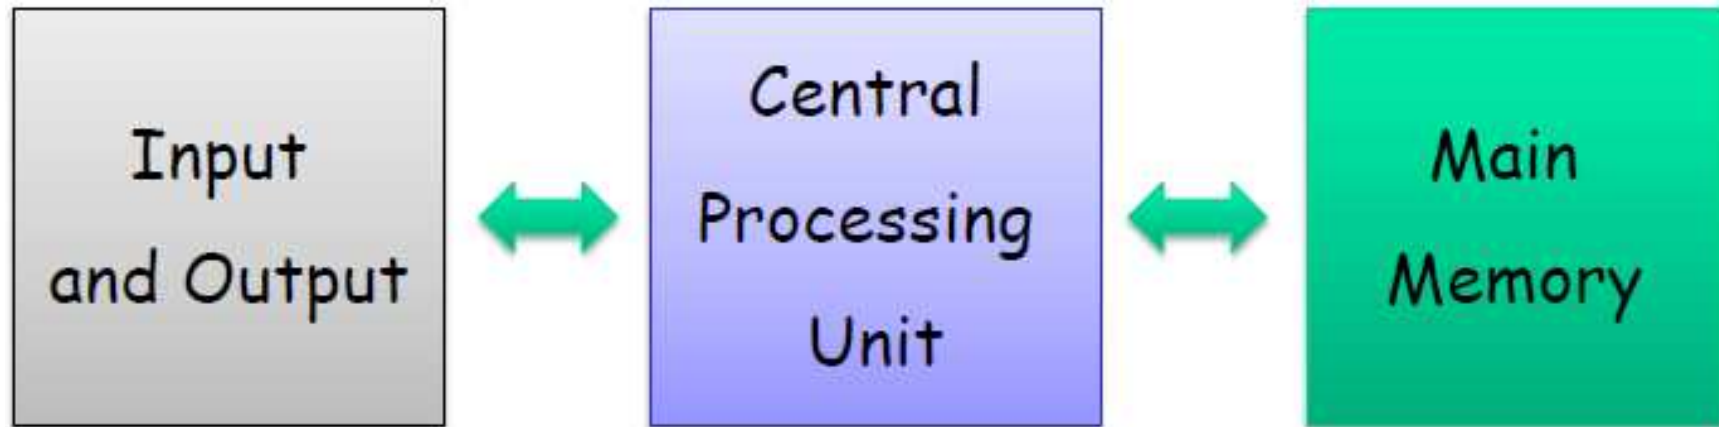
- ❑ PC architecture
- ❑ x86 Instruction set
- ❑ gcc calling conventions

PC Organization

- ❑ One or more CPUs, memory, and device controllers connect through system bus

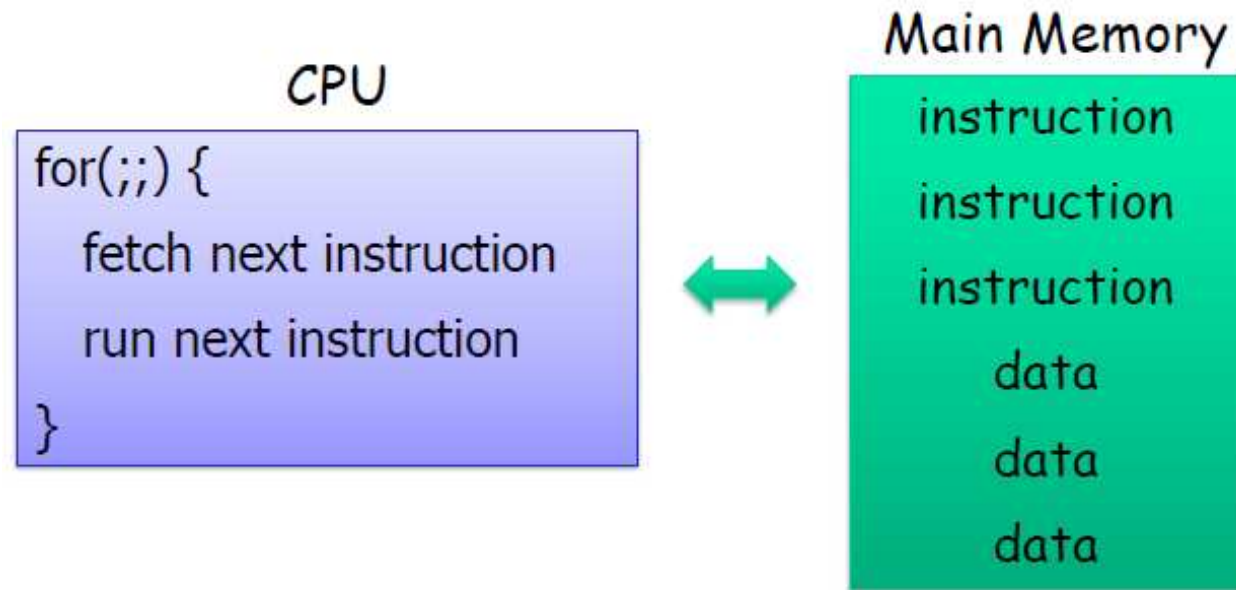


Abstract Model



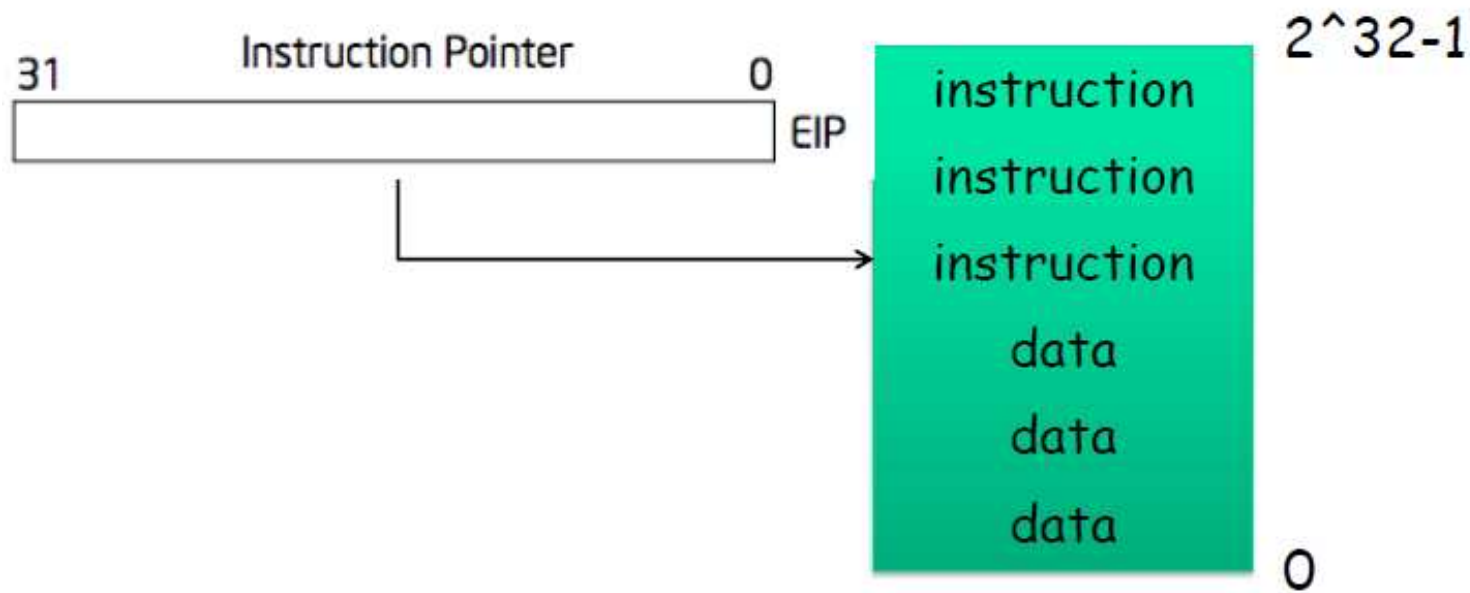
- ❑ I/O: communicating data to and from devices
- ❑ CPU: digital logic for doing computation
- ❑ Memory: N words of B bits

The Stored Program Computer



- ❑ Memory holds both instructions and data
- ❑ CPU interprets instructions
- ❑ Instructions read/write data

x86 Implementation



- ❑ **EIP** incremented after each instruction
- ❑ Variable length instructions
- ❑ **EIP** modified by **CALL**, **RET**, **JMP**,
conditional **JMP**

Registers: Work Space

General-Purpose Registers					
31	16 15	8 7	0	16-bit	32-bit
	AH	AL		AX	EAX
	BH	BL		BX	EBX
	CH	CL		CX	ECX
	DH	DL		DX	EDX
	BP				EBP
	SI				ESI
	DI				EDI
	SP				ESP

ESP: stack pointer

EBP: frame base pointer

ESI: source index

EDI: destination index


❑ 8, 16, and 32 bit versions

❑ Example: **ADD EAX, 10**

❑ More: **SUB, AND**, etc.

❑ By convention some for special purposes

EFLAGS Register



❑ Track current CPU status

x86 Instruction Set

❑ Instruction syntax

- ❑ Intel manual Volume 2: `op dst, src`
- ❑ AT&T (gcc/gas): `op src, dst`
 - ❑ `op` uses suffix b, w, l for 8, 16, 32-bit operands
 - ❑ xv6, JOS

❑ Instruction classes

- ❑ Data movement: `MOV, PUSH, POP, ...`
- ❑ Arithmetic: `TEST, SHL, ADD, AND, ...`
- ❑ I/O: `IN, OUT, ...`
- ❑ Control: `JMP, JZ, JNZ, CALL, RET`
- ❑ String: `MOVS, REP, ...`
- ❑ System: `INT, IRET`

Memory: Addressing Examples

Address	Value	Operand	Value	Comment
0x100	0xFF	%eax	0x100	Value in the register
0x104	0xAB	0x104	0xAB	Value is at the address
0x108	0x13	\$0x108	0x108	Value is the value (\$ means immediate, i.e. constant, value)
0x10C	0x11	(%eax)	0xFF	value is at the address stored in the register -> GTV@(reg)
		4(%eax)	0xAB	GTV@(4 + reg)
Register	Value	9(%eax, %edx)	0x11	GTV@(9 + reg + reg)
%eax	0x100	0xFC(, %ecx, 4)	0xFF	GTV@(0xFC + 0 + reg*4)
%ecx	0x1	(%eax, %edx, 4)	0x11	GTV@(reg + reg*4)
%edx	0x3			

❑ Red means memory address

Memory: Addressing Modes

<code>movl %eax, %edx</code>	<code>edx = eax;</code>	<i>register mode</i>
<code>movl \$0x123, %edx</code>	<code>edx = 0x123;</code>	<i>immediate</i>
<code>movl 0x123, %edx</code>	<code>edx = *(int32_t*)0x123;</code>	<i>direct</i>
<code>movl (%ebx), %edx</code>	<code>edx = *(int32_t*)ebx;</code>	<i>indirect</i>
<code>movl 4(%ebx), %edx</code>	<code>edx = *(int32_t*)(ebx+4);</code>	<i>displaced</i>

- ❑ Memory instructions: **MOV**, **PUSH**, **POP**, etc.
- ❑ More instructions can take a memory address

Stack Memory + Operations

<u>Example instruction</u>	<u>What it does</u>
<code>pushl %eax</code>	<code>subl \$4, %esp</code> <code>movl %eax, (%esp)</code>
<code>popl %eax</code>	<code>movl (%esp), %eax</code> <code>addl \$4, %esp</code>
<code>call 0x12345</code>	<code>pushl %eip (*)</code> <code>movl \$0x12345, %eip (*)</code>
<code>ret</code>	<code>popl %eip (*)</code>

❑ For implementing function calls

❑ Stack grows “down” on x86


Stack Operation Example

`pushl $8` ①

`movl %esp, %ebp` ②

`subl $4, %esp` ③

`movl $8, (%esp)` ④



More Memory

- ❑ 8086 16-bit register and 20-bit bus addresses
- ❑ These extra 4 bits come from segment register
 - ❑ **CS**: code segment, for EIP
 - ❑ Instruction address: $CS * 16 + EIP$
 - ❑ **SS**: stack segment, for **ESP** and **EBP**
 - ❑ **DS**: data segment for load/store via other registers
 - ❑ **ES**: another data segment, destination for string ops
- ❑ Make life more complicated
 - ❑ Cannot directly use 16-bit stack address as pointer
 - ❑ For a far pointer programmer must specify segment reg
 - ❑ Pointer arithmetic and array indexing across seg bound

And More Memory

- ❑ 80386: 32 bit register and addresses (1985)
- ❑ AMD k8: 64 bit (2003)
 - ❑ **RAX** instead of **EAX**
 - ❑ x86-64, x64, amd64, intel64: all same thing
- ❑ Backward compatibility
 - ❑ Boots in 16-bit mode; **bootasm.S** switches to 32
 - ❑ Prefix **0x66** gets 32-bit mode instructions
 - ❑ **MOVW** in 32-bit mode = **0x66** + **MOVW** in 16-bit mode
 - ❑ **.code32** in **bootasm.S** tells assembler to insert **0x66**
- ❑ 80386 also added virtual memory addresses

I/O Space and Instructions

```
#define DATA_PORT    0x378
#define STATUS_PORT   0x379
#define    BUSY    0x80
#define CONTROL_PORT  0x37A
#define    STROBE   0x01
void
lpt_putc(int c)
{
    /* wait for printer to consume previous byte */
    while((inb(STATUS_PORT) & BUSY) == 0)
        ;

    /* put the byte on the parallel lines */
    outb(DATA_PORT, c);

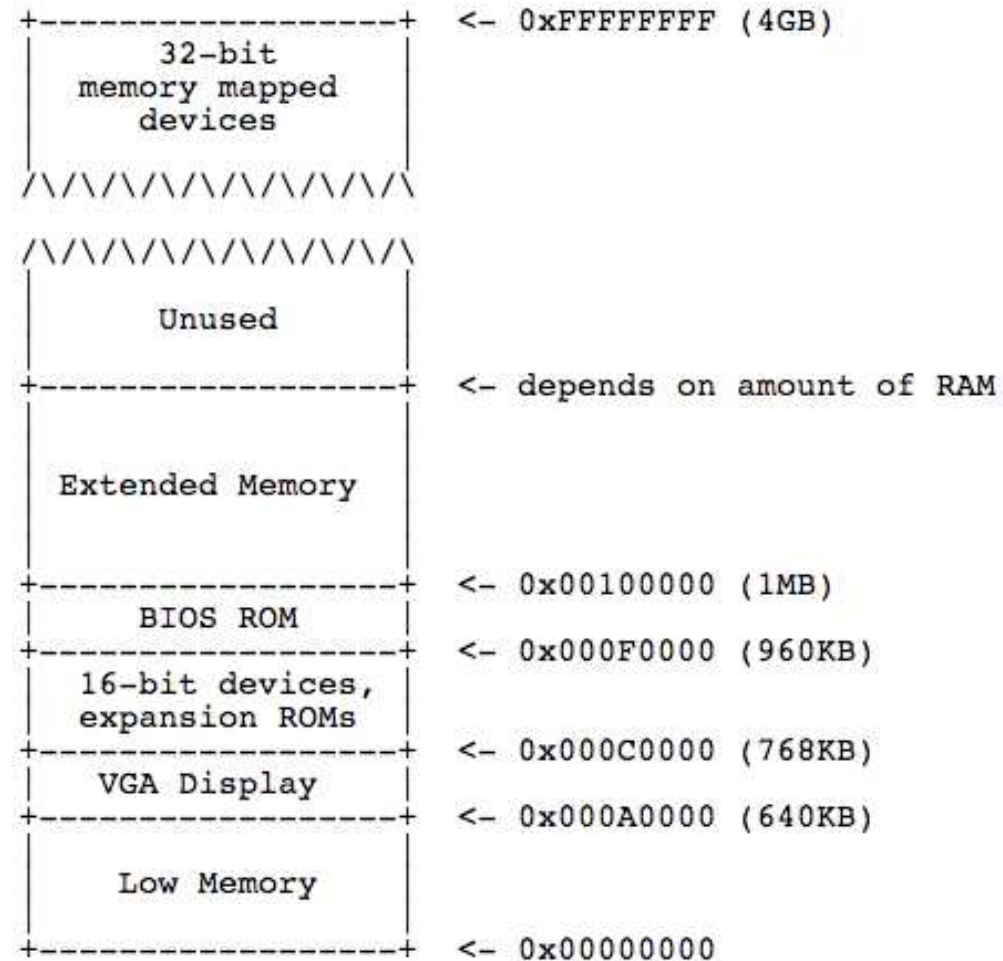
    /* tell the printer to look at the data */
    outb(CONTROL_PORT, STROBE);
    outb(CONTROL_PORT, 0);
}
```

❑ 8086: only 1024 I/O addresses

Memory-mapped I/O

- ❑ Use normal addresses for I/O
 - ❑ No special instructions
 - ❑ No 1024 limit
 - ❑ Hardware routes to appropriate device
- ❑ Works like “magic” memory
 - ❑ I/O device addressed and accessed like memory
 - ❑ However, reads and writes have “side effects”
 - ❑ Read result can change due to external events

Memory Layout



gcc Inline Assembly

- ❑ Can embed assembly code in C code

 - ❑ Many examples in xv6

- ❑ Basic syntax: `asm (“assembly code”)`

 - ❑ e.g., `asm (“movl %eax, %ebx”)`

- ❑ Advanced syntax:

 - `asm (assembler template`

 - `: output operands /* optional */`

 - `: input operands /* optional */`

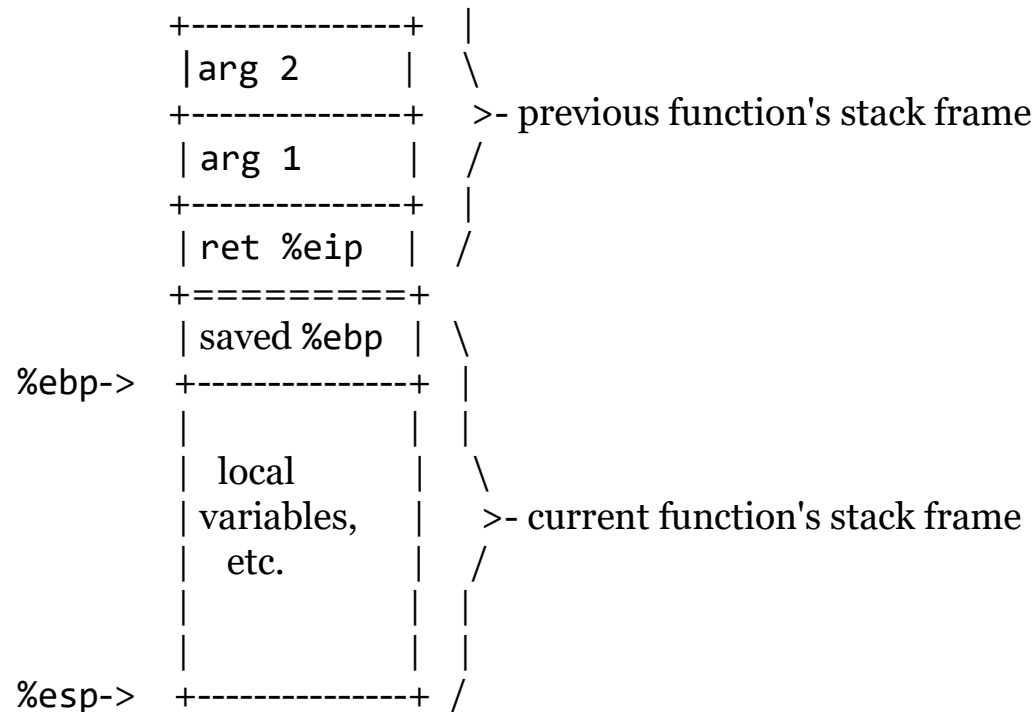
 - `: list of clobbered registers /* optional */);`

 - e.g., `int val;`

 - `asm (“movl %%ebp,%0” : “=r” (val));`

gcc Calling Conventions

- ❑ Functions can do anything that doesn't violate contract. By convention, GCC does more:
 - ❑ each function has a stack frame marked by `%ebp`, `%esp`
 - ❑ `%esp` can move to make stack frame bigger, smaller
 - ❑ `%ebp` points at saved `%ebp` from previous function, chain to walk stack



gcc Calling Conventions (cont.)

- ❑ `%eax` contains return value, `%ecx`, `%edx` may be trashed
- ❑ 64 bit return value: `%eax + %edx`
- ❑ `%ebp`, `%ebx`, `%esi`, `%edi` must be as before call
- ❑ Caller saved: `%eax`, `%ecx`, `%edx`
- ❑ Callee saved: `%ebp`, `%ebx`, `%esi`, `%edi`