

# CPU Scheduling (2)

---

Dr. Jun Zheng

CSE325 Principles of Operating  
Systems


9/11/2019



# Determining Length of Next CPU Burst

- ❑ Can only estimate the length – should be similar to the previous one
  - ❑ Then pick process with shortest predicted next CPU burst
    1.  $t_n$  = actual length of  $n^{th}$  CPU burst
    2.  $\tau_{n+1}$  = predicted value for the next CPU burst
    3.  $\alpha, 0 \leq \alpha \leq 1$
    4. Define:  $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$ .
- ❑ Can be done by using the length of previous CPU bursts, using exponential averaging
- ❑ Commonly,  $\alpha$  set to  $1/2$
- ❑ Preemptive version called **shortest-remaining-time-first**

# Prediction of the Length of the Next CPU Burst



CPU burst ( $t_i$ )	6	4	6	4	13	13	13	...
"guess" ( $\tau_i$ )	10	8	6	6	9	11	12	...

# Examples of Exponential Averaging

□  $\alpha = 0$

□  $\tau_{n+1} = \tau_n$

□ Recent history does not count

□  $\alpha = 1$

□  $\tau_{n+1} = \alpha t_n$

□ Only the actual last CPU burst counts

□ If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$


□ Since both  $\alpha$  and  $(1 - \alpha)$  are less than or equal to 1, each successive term has less weight than its predecessor

# Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

- *Preemptive* SJF Gantt Chart



- Average waiting time =  $[(10-1)+(1-1)+(17-2)+(5-3)]/4 = 26/4 = 6.5$  msec

# Priority Scheduling

- ❑ A priority number (integer) is associated with each process
- ❑ The CPU is allocated to the process with the highest priority
  - ❑ Preemptive
  - ❑ Nonpreemptive
- ❑ SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- ❑ Problem = **Starvation** – low priority processes may never execute
- ❑ Solution = **Aging** – as time progresses increase the priority of the process

# Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

- Priority scheduling Gantt Chart



- Average waiting time = 8.2 msec

# Round Robin (RR)


- ❑ Each process gets a small unit of CPU time (**time quantum**  $q$ ), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ❑ If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units.
- ❑ Timer interrupts every quantum to schedule next process
- ❑ Performance
  - ❑  $q$  large  $\Rightarrow$  FIFO
  - ❑  $q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high



# Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- The Gantt chart is:



- Typically, higher average turnaround than SJF, but better **response**
- $q$  should be large compared to context switch time
- $q$  usually 10ms to 100ms, context switch < 10 usec

# Average Waiting Time

$$P_1: (0-0) + (10 - 4) = 6$$

$$P_2: (4-0) = 4$$

$$P_3: (7 - 0) = 7$$

$$\text{Avg. waiting time: } (6 + 4 + 7)/3 = 17/3 \text{ ms}$$