

# Threads

---

Dr. Jun Zheng

CSE325 Principles of Operating  
Systems

9/18/2019



# Processes

- ❑ A process includes many things
  - ❑ An address space (defining all the code and data pages)
  - ❑ OS resources (e.g., open files) and accounting information
  - ❑ Execution state (PC, SP, regs, etc.)
- ❑ Creating a new process is costly because of all of the data structures that must be allocated and initialized
  - ❑ Recall `struct task_struct` in Linux
  - ❑ ...which does not even include page tables, perhaps TLB flushing, etc.
- ❑ Communicating between processes is costly because most communication goes through the OS
  - ❑ Overhead of system calls and copying data

# Multi-programming

- ❑ To execute parallel programs we need to
  - ❑ Create several processes that execute in parallel
  - ❑ Cause each to map to the same address space to share data
    - ❑ They are all part of the same computation
  - ❑ Have the OS schedule these processes in parallel
- ❑ This situation is **very inefficient**
  - ❑ **Space**: PCB, page tables, etc.
  - ❑ **Time**: create data structures, fork and copy address space, etc.
- ❑ Solutions: possible to have more **efficient**, yet **cooperative** “processes”?

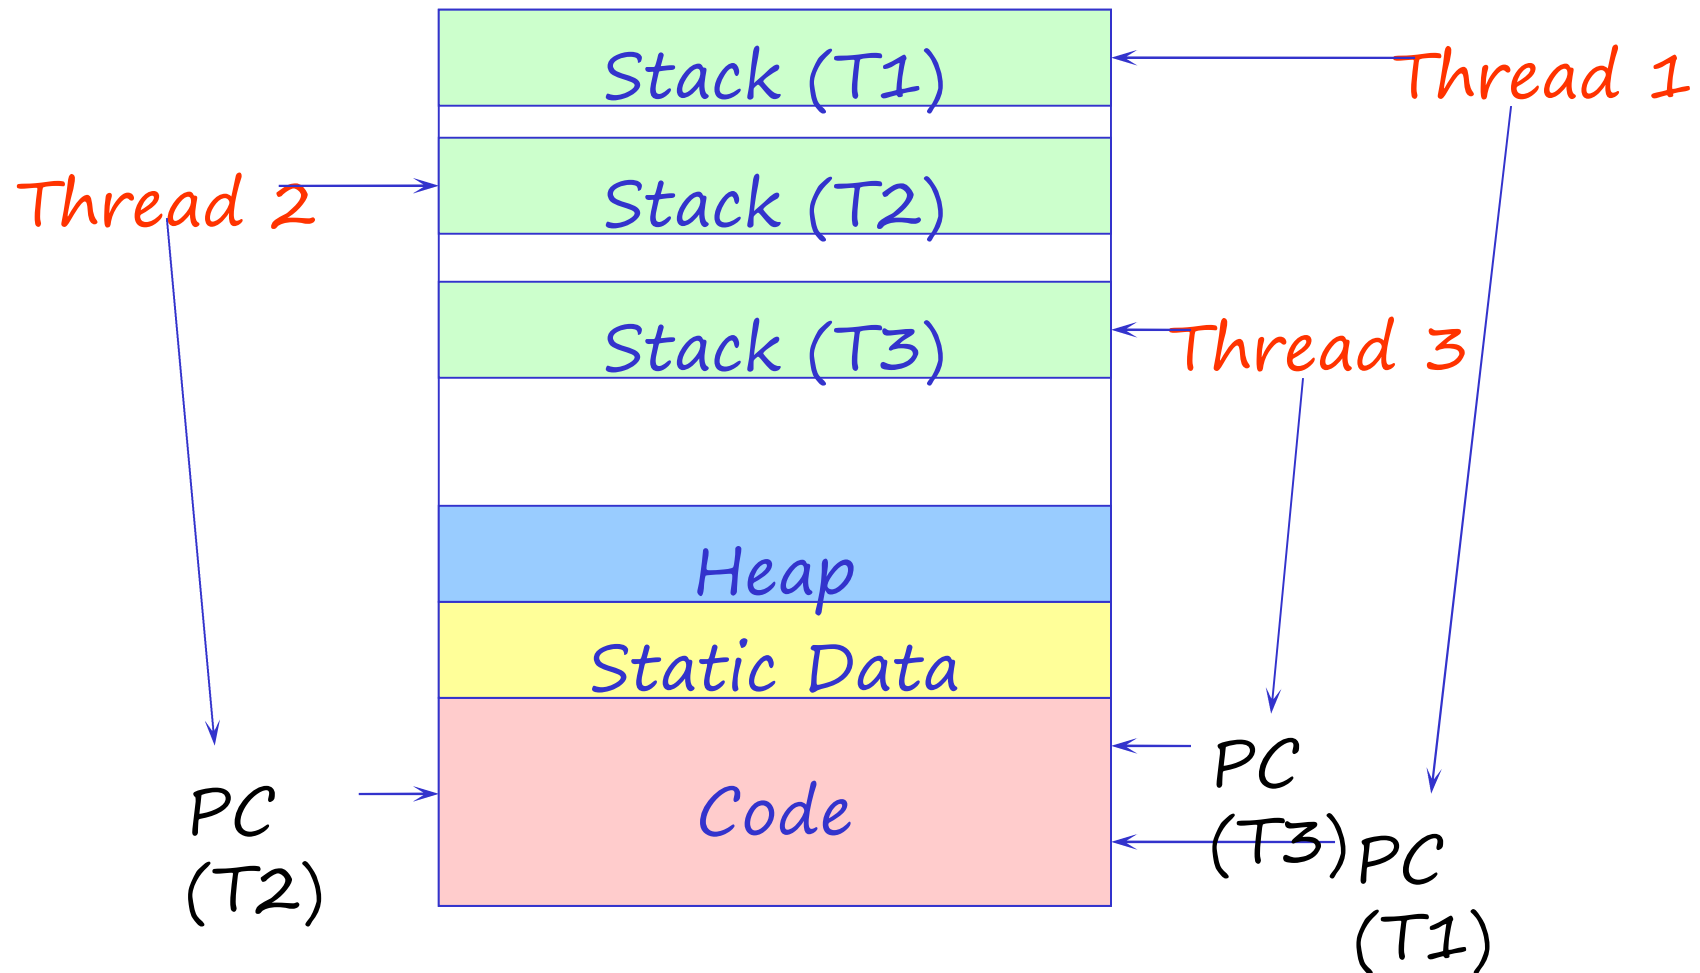
# Rethinking Processes

- ❑ What is similar in these cooperating processes?
  - ❑ They all share the same code and data (address space)
  - ❑ They all share the same privileges
  - ❑ They all share the same resources (files, sockets, etc.)
- ❑ What don't they share?
  - ❑ Each has its own execution state: PC, SP, and registers
- ❑ **Key idea:** Why don't we separate the concept of a process from its execution state?
  - ❑ **Process:** address space, privileges, resources, etc.
  - ❑ **Execution state:** PC, SP, registers
- ❑ Exec state also called **thread of control**, or **thread**

# Threads

- ❑ Modern OSes (Mac, Windows, modern Unix) separate the concepts of processes and threads
  - ❑ The **thread** defines a sequential execution stream within a process (PC, SP, registers)
  - ❑ The **process** defines the address space and general process attributes (everything but threads of execution)
- ❑ A thread is bound to a single process
  - ❑ Processes, however, can have **multiple** threads
- ❑ Threads become the unit of scheduling
  - ❑ Processes are now the **containers** in which threads execute

# Threads in a Process



# Thread Design Space

