

Threads (2)

Dr. Jun Zheng
CSE325 Principles of Operating
Systems
9/20/2019



Threads: Concurrent Servers

- ❑ Using **fork()** to create new processes to handle requests in parallel is overkill for such a simple task
- ❑ Web server example:



```
- while (1) {  
-     int sock = accept();  
-     if ((child_pid = fork()) == 0) {  
-         Handle client request  
-         Close socket and exit  
-     } else {  
-         Close socket  
-     }  
- }
```

Threads: Concurrent Servers

- ❑ Instead, we can create a new thread for each request

```
⑩ web_server() {  
  - while (1) {  
    - int sock = accept();  
    - thread_create(handle_request, sock);  
  - }  
⑩ }  
⑩ handle_request(int sock) {  
⑩     Process request  
⑩     close(sock);  
⑩ }
```

Concurrency vs. Parallelism



Benefits of Multithreaded Programming

- ❑ **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- ❑ **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- ❑ **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- ❑ **Scalability** – process can take advantage of multiprocessor architectures

Question

- ❑ If a process exits and there are still threads of that process running, will they continue to run? Please justify your answer.

Answer

No. When a process exits, it takes everything with it—the process structure, the memory space, everything—including threads.

Kernel-Level Threads

- ❑ We have taken the execution aspect of a process and separated it out into threads
 - ❑ To make concurrency cheaper
- ❑ As such, the OS now manages threads *and* processes
 - ❑ All thread operations are implemented in the kernel
 - ❑ The OS schedules all of the threads in the system
- ❑ OS-managed threads are called **kernel-level threads** or **lightweight processes**
 - ❑ Windows: **threads**
 - ❑ Solaris: **lightweight processes (LWP)**
 - ❑ POSIX Threads (pthreads):
PTHREAD_SCOPE_SYSTEM

Kernel-level Thread Limitations

- ❑ Kernel-level threads make concurrency much cheaper than processes
 - ❑ Much less state to allocate and initialize
- ❑ However, for fine-grained concurrency, kernel-level threads still suffer from too much overhead
 - ❑ Thread operations still require system calls
 - ❑ Ideally, want thread operations to be **as fast as a procedure call**
- ❑ For such fine-grained concurrency, need even “cheaper” threads

User-Level Threads

- ❑ To make threads cheap and fast, they need to be implemented at user level
 - ❑ Kernel-level threads are managed by the OS
 - ❑ User-level threads are managed entirely by the run-time system (user-level library)
- ❑ User-level threads are small and fast
 - ❑ A thread is simply represented by a PC, registers, stack, and small thread control block (TCB)
 - ❑ Creating a new thread, switching between threads, and synchronizing threads are done via procedure call
 - ❑ No kernel involvement
 - ❑ User-level thread operations **100x faster** than kernel threads
 - ❑ pthreads: **PTHREAD_SCOPE_PROCESS**

User-level Thread Limitations

- ❑ But, user-level threads are not a perfect solution
 - ❑ As with everything else, they are a tradeoff
- ❑ User-level threads are **invisible** to the OS
 - ❑ They are not well integrated with the OS
- ❑ As a result, the OS can make poor decisions
 - ❑ Scheduling a process with idle threads
 - ❑ **Blocking a process whose thread initiated an I/O, even though the process has other threads that can execute**
- ❑ Solving this requires communication between the kernel and the user-level thread manager

Kernel- vs. User-level Threads

- ❑ Kernel-level threads
 - ❑ Integrated with OS (informed scheduling)
 - ❑ Slow to create, manipulate, synchronize
- ❑ User-level threads
 - ❑ Fast to create, manipulate, synchronize
 - ❑ Not integrated with OS (uninformed scheduling)
- ❑ Understanding the differences between kernel- and user-level threads is important

Question

- ❑ A disadvantage of ULTs is that when a ULT executes a blocking system call (e.g. a I/O call), not only is that thread blocked, but all of the threads within the process are blocked. Why?

Answer


- ❑ Because, with ULTs, the thread structure of a process is not visible to the operating system, which only schedules on the basis of processes.
- ❑ For a blocking system call like a I/O call, the control is transferred to the kernel. The kernel invokes the I/O action, places process in the blocked state, and switches to another process.

Multithreading Models

- ❑ Many-to-One
- ❑ One-to-One
- ❑ Many-to-Many


Many-to-One

- ❑ Many user-level threads mapped to single kernel thread
- ❑ One thread blocking causes all to block
- ❑ Multiple threads may not run in parallel on muticore system because only one may be in kernel at a time
- ❑ Few systems currently use this model
- ❑ Examples:
 - ❑ **Solaris Green Threads**
 - ❑ **GNU Portable Threads**




One-to-One

- ❑ Each user-level thread maps to kernel thread
- ❑ Creating a user-level thread creates a kernel thread
- ❑ More concurrency than many-to-one
- ❑ Number of threads per process sometimes restricted due to overhead
- ❑ Examples
 - ❑ Windows
 - ❑ Linux
 - ❑ Solaris 9 and later



Many-to-Many Model

- ❑ Allows many user level threads to be mapped to many kernel threads
- ❑ Allows the operating system to create a sufficient number of kernel threads
- ❑ Solaris prior to version 9
- ❑ Windows with the *ThreadFiber* package
- ❑ Extremely difficult to implement



Two-level Model

- ❑ Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- ❑ Examples
 - ❑ IRIX
 - ❑ HP-UX
 - ❑ Tru64 UNIX
 - ❑ Solaris 8 and earlier

