

# Deadlocks (2)

---

Dr. Jun Zheng

CSE325 Principles of Operating  
Systems

10/9/2019



# System Model

- ❑ System consists of resources
- ❑ Resource types  $R_1, R_2, \dots, R_m$ 
  - ❑ *CPU cycles, memory space, I/O devices*
- ❑ Each resource type  $R_i$  has  $W_i$  instances.
- ❑ Each process utilizes a resource as follows:
  - ❑ **request**
  - ❑ **use**
  - ❑ **release**

# Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- ❑ **Mutual exclusion:** only one process at a time can use a resource
- ❑ **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- ❑ **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task
- ❑ **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .


# Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .


- $V$  is partitioned into two types:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
- **request edge** – directed edge  $P_i \rightarrow R_j$
- **assignment edge** – directed edge  $R_j \rightarrow P_i$

# Resource-Allocation Graph (Cont.)


□ Process




□ Resource Type with 4 instances




□  $P_i$  requests instance of  $R_j$




□  $P_i$  is holding an instance of  $R_j$




# Example of a Resource Allocation Graph



# Resource Allocation Graph With A Deadlock



# Graph With A Cycle But No Deadlock





# Basic Facts

- ❑ If graph contains no cycles  $\Rightarrow$  no deadlock
- ❑ If graph contains a cycle  $\Rightarrow$ 
  - ❑ if only one instance per resource type, then deadlock
  - ❑ if several instances per resource type, possibility of deadlock

# Methods for Handling Deadlocks

- ❑ Ensure that the system will *never* enter a deadlock state:
  - ❑ Deadlock prevention
  - ❑ Deadlock avoidance
- ❑ Allow the system to enter a deadlock state and then detect and recover
- ❑ Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX

# Deadlock Prevention

Restrain the ways request can be made

- ❑ **Mutual Exclusion** – not required for sharable resources (e.g., read-only files); must hold for non-sharable resources
- ❑ **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
  - ❑ Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none allocated to it.
  - ❑ Low resource utilization; starvation possible

# Deadlock Prevention (Cont.)

## ☐ **No Preemption** –

- ☐ If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
- ☐ Preempted resources are added to the list of resources for which the process is waiting
- ☐ Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting

## ☐ **Circular Wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

# Deadlock Avoidance

Requires that the system has some additional *a priori* information available

- ❑ Simplest and most useful model requires that each process declare the ***maximum number*** of resources of each type that it may need
- ❑ The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- ❑ Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes


# Safe State

- ❑ When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- ❑ System is in **safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the systems such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$
- ❑ That is:
  - ❑ If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - ❑ When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - ❑ When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on

# Basic Facts

- ❑ If a system is in safe state  $\Rightarrow$  no deadlocks
- ❑ If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- ❑ Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.

# Safe, Unsafe, Deadlock State





# Avoidance Algorithms

- ❑ Single instance of a resource type
  - ❑ Use a resource-allocation graph
  
- ❑ Multiple instances of a resource type
  - ❑ Use the banker's algorithm