


Memory Management (4)

Dr. Jun Zheng
CSE325 Principles of Operating
Systems

10/25/2019




Page Protection Example



Virtual
Memory

		pwu
0	1	101
1	4	110
2	3	000
3	7	111

Page table




Physical
Memory

Page Allocation

- ❑ Free frame management
 - ❑ E.g., can put page on a **free list**
- ❑ Allocation policy
 - ❑ E.g., one page at a time, from head of free list
- ❑ xv6: **kalloc.c**

Page Allocation



Before allocation

After allocation

Implementation of Page Table


- ❑ Page table is kept in main memory
 - ❑ **Page-table base register (PTBR)** points to the page table
 - ❑ x86: **cr3**
 - ❑ **Page-table length register (PTLR)** indicates size of the page table
 - ❑ OS stores base in process control block (PCB)
 - ❑ OS switches PTBR on each context switch
- ❑ In this scheme every data/instruction access requires two memory accesses
 - ❑ One for the page table and one for the data / instruction
- ❑ The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**

Structure of the Page Table

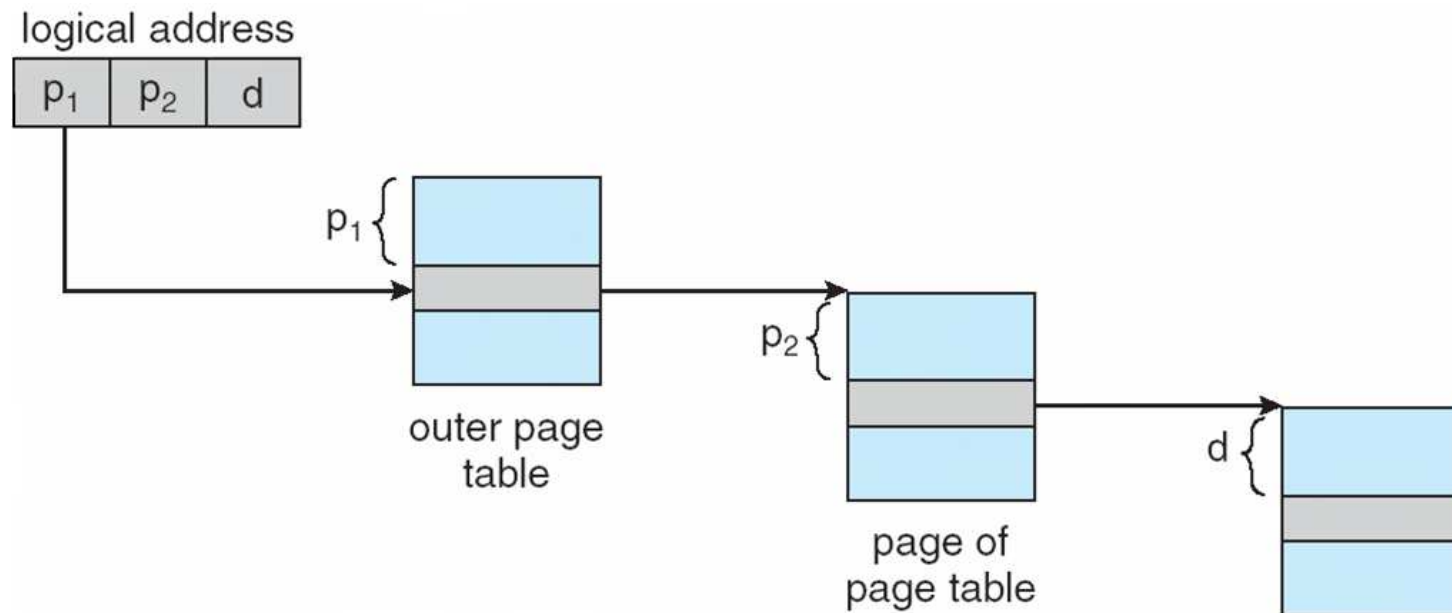
- ❑ Memory structures for paging can get huge using straight-forward methods
 - ❑ Consider a 32-bit logical address space as on modern computers
 - ❑ Page size of 4 KB (2^{12})
 - ❑ Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - ❑ If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
 - ❑ That amount of memory used to cost a lot
 - ❑ Don't want to allocate that contiguously in main memory
- ❑ Hierarchical Paging
- ❑ Hashed Page Tables
- ❑ Inverted Page Tables

Hierarchical Page Tables

- ❑ Break up the logical address space into multiple page tables
- ❑ A simple technique is a two-level page table
- ❑ We then page the page table




Address Translation with Hierarchical Page Table

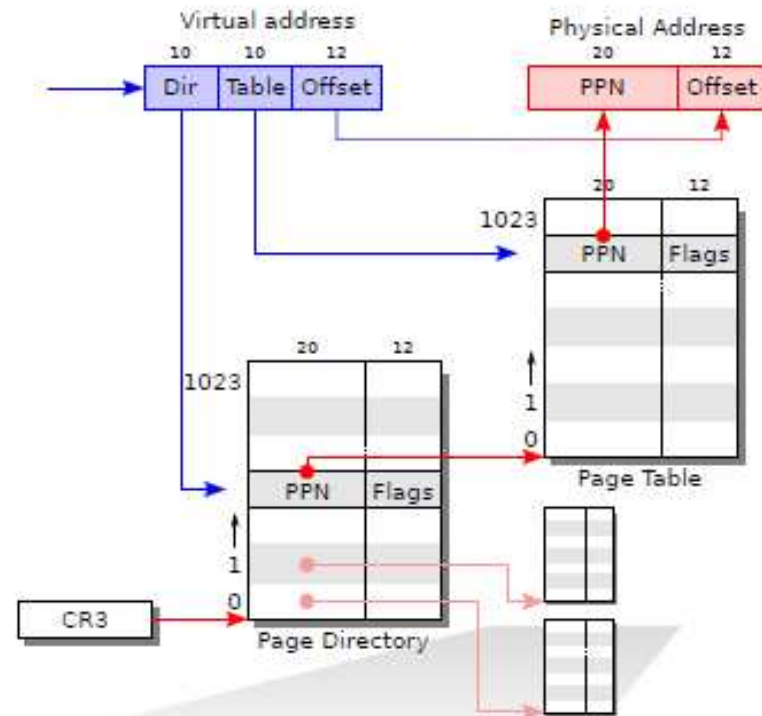
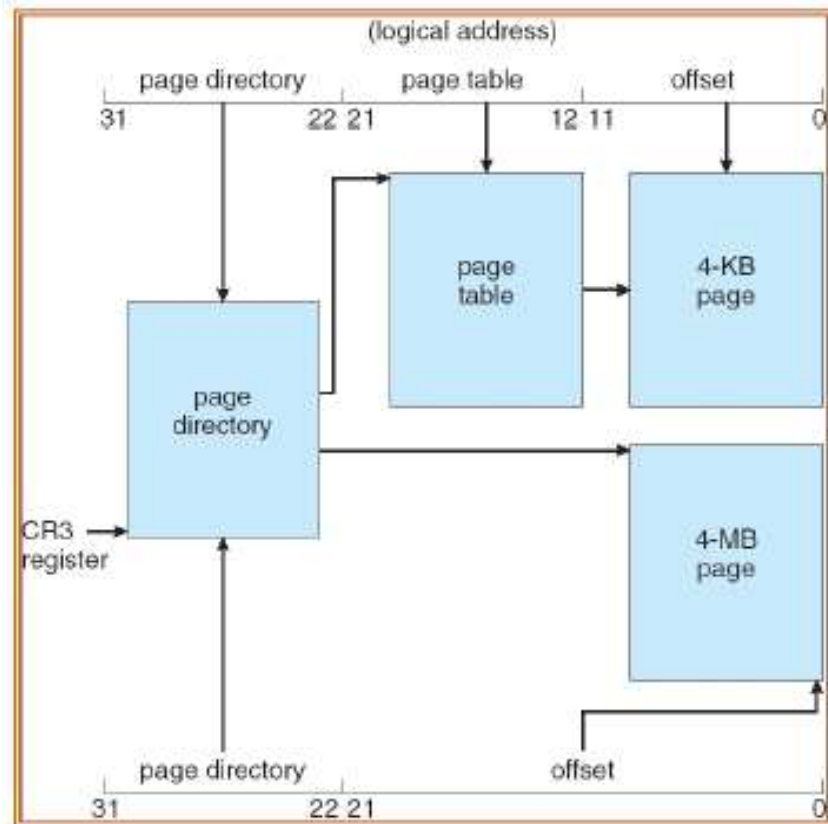


x86 Page Translation

- ❑ 32-bit address space, 4 KB page
 - ❑ 4KB page -> 12 bits for page offset
- ❑ How many bits for 2nd-level page table?
 - ❑ Desirable to fit a 2nd-level page table in one page
 - ❑ $4\text{KB}/4\text{B} = 1024 \rightarrow 10$ bits for 2nd-level page table
- ❑ Address bits for top-level page table: $32 - 12 - 10 = 10$




x86 Paging Architecture



64-bit Logical Address Space

- ❑ Even two-level paging scheme not sufficient
- ❑ If page size is 4 KB (2^{12})
 - ❑ Then page table has 2^{52} entries
 - ❑ If two level scheme, inner page tables could be 2^{10} 4-byte entries
 - ❑ Address would look like




- ❑ Outer page table has 2^{42} entries or 2^{44} bytes
- ❑ One solution is to add a 2nd outer page table
- ❑ But in the following example the 2nd outer page table is still 2^{34} bytes in size
 - ❑ And possibly 4 memory access to get to one physical memory location

Hashed Page Tables

- ❑ Common in address spaces > 32 bits
- ❑ The logical page number is hashed into a page table
 - ❑ This page table contains a chain of elements hashing to the same location
- ❑ Each element contains (1) the logical page number (2) the value of the mapped page frame (3) a pointer to the next element
- ❑ Logical page numbers are compared in this chain searching for a match
 - ❑ If a match is found, the corresponding physical frame is extracted
- ❑ Variation for 64-bit addresses is **clustered page tables**
 - ❑ Similar to hashed but each entry refers to several pages (such as 16) rather than 1
 - ❑ Especially useful for **sparse** address spaces (where memory references are non-contiguous and scattered)

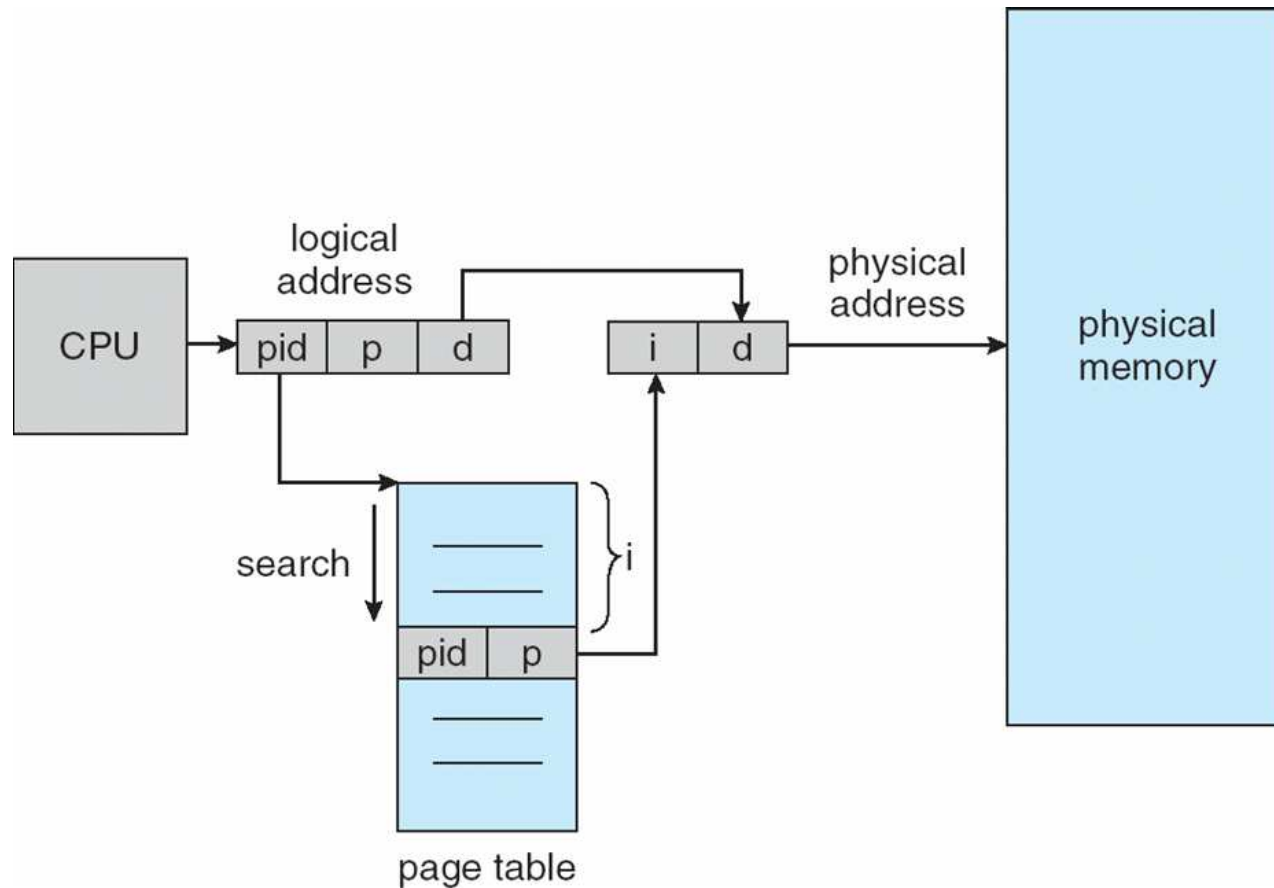
Hashed Page Table



Inverted Page Table

- ❑ Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- ❑ One entry for each real page of memory
- ❑ Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- ❑ Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- ❑ Use hash table to limit the search to one — or at most a few — page-table entries
 - ❑ TLB can accelerate access
- ❑ But how to implement shared memory?
 - ❑ One mapping of a logical address to the shared physical address

Inverted Page Table Architecture




Example

The VAX architecture supports a variation of two-level paging.

The VAX is a 32-bit machine with a page size of 512 bytes.

The logical address space of a process is divided into four equal sections, each of which consists of 2^{30} bytes. Each section represents a different part of the logical address space of a process. The first 2 high-order bits of the logical address designate the appropriate section. The next 21 bits represent the logical page number of that section, and the final 9 bits represent an offset in the desired page.



How many memory operations are performed when a user program executes a memory-load operation?

Example

Consider a computer system with a 32-bit logical address and 4-KB (2^{12} B) page size. The system supports up *to* 512MB (2^{29} B) of physical memory. How many entries are there in each of the following?

- (1) A conventional single-level page table
- (2) An inverted page table

Answer

(1) # of pages = # of entries = ????

Size of logical address space = 2^m = # of pages \times page size

»» 2^{32} = # of pages $\times 2^{12}$

of pages = $2^{32} / 2^{12} = 2^{20}$ pages

(2)

Size of physical address space = # of frames \times frame size

(frame size = page size)

For inverted page table, # of frames = # of entries

2^{29} = # of frames $\times 2^{12}$

of frames = $2^{29} / 2^{12} = 2^{17}$

of entries = 2^{17}