

Virtual Memory (2)

Dr. Jun Zheng
CSE325 Principles of Operating
Systems
11/6/2019



Performance of Demand Paging

Page Fault Rate

- $0 \leq p \leq 1.0$
- if $p = 0$ no page faults
- if $p = 1$, every reference is a fault

Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & \quad + \text{swap page out} \\ & \quad + \text{swap page in} \\ & \quad + \text{restart overhead}) \end{aligned}$$


Demand Paging Example

- ❑ Memory access time = 200ns
- ❑ Average page fault service time = 8ms
- ❑ $EAT = (1-p)*200 + p*8,000,000 = 200 + p*7,999,800$
- ❑ If one access out of 1000 causes a page fault, then $EAT = 8.2\mu s$
 - ❑ A slowdown by a factor of 40

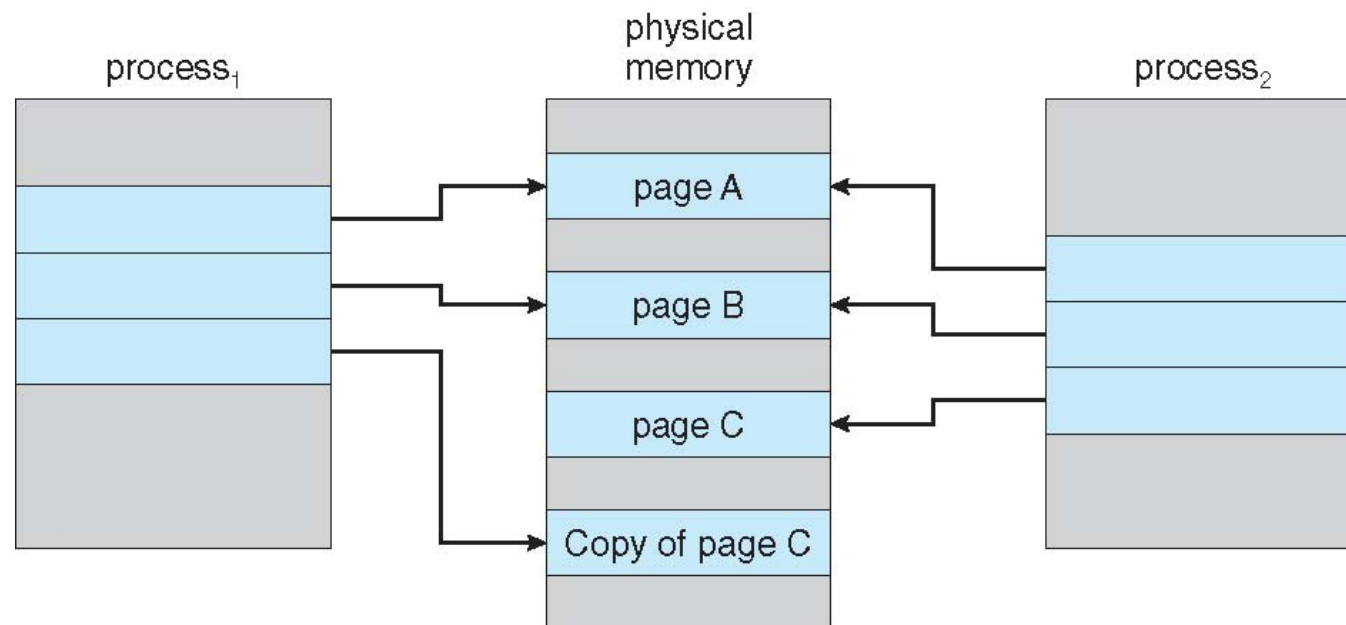
Copy-on-Write

- ❑ **Copy-on-Write** (COW) allows both parent and child processes to initially *share* the same pages in memory
 - ❑ If either process modifies a shared page, only then is the page copied
- ❑ COW allows more efficient process creation as only modified pages are copied
- ❑ **vfork()** variation on **fork()** system call has parent suspend and child using copy-on-write address space of parent
 - ❑ Designed to have child call **exec()**
 - ❑ Very efficient

Before Process 1 Modifies Page C



After Process 1 Modifies Page C



What Happens if There is no Free Frame?

- ❑ Used up by process pages
- ❑ Also in demand from the kernel, I/O buffers, etc
- ❑ How much to allocate to each?
- ❑ Page replacement – find some page in memory, but not really in use, page it out
 - ❑ Performance – want an algorithm which will result in minimum number of page faults

Page Replacement

- ❑ Prevent **over-allocation** of memory by modifying page-fault service routine to include page replacement
- ❑ Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- ❑ Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory

Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim frame**
 - Write victim frame to disk if dirty
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Continue the process by restarting the instruction that caused the trap

Page Replacement Algorithms

- ❑ Want lowest page-fault rate on both first access and re-access
- ❑ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
 - ❑ String is just page numbers, not full addresses
 - ❑ Repeated access to the same page does not cause a page fault
 - ❑ Results depend on number of frames available
- ❑ In the examples if not specified, the **reference string** of referenced page numbers is


7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1

Random Replacement

- ❑ Replaced page is chosen from m loaded frames with probability $1/m$
- ❑ Easy to implement but does not perform well

Optimal Algorithm


- ❑ Replace page that will not be used for longest period of time



- ❑ How do you know this?
 - ❑ Can't read the future
- ❑ Used for measuring how well your algorithm performs

First-In-First-Out (FIFO) Algorithm

- ❑ Reference string:
7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1
- ❑ 3 frames (3 pages can be in memory at a time per process)



15 page faults

- ❑ What if we have 4 frames?
- ❑ Consider reference string: 1 2 3 4 1 2 5 1 2 3 4 5

Belady's Anomaly

❑ Adding more frames can cause more page faults!

