

# Virtual Memory (4)

---

Dr. Jun Zheng  
CSE325 Principles of Operating  
Systems  
11/11/2019



# In-Class Work 6

Consider the following page reference string:

1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3,  
6

How many page faults would occur for the following replacement algorithms, assuming **three frames**?

Remember that all frames are initially empty, so your first unique pages will cost one fault each.

- ☐ LRU replacement
- ☐ FIFO replacement
- ☐ Optimal replacement

# LRU Approximation Algorithms


- ❑ LRU needs special hardware and still slow
- ❑ **Reference bit**
  - ❑ With each page associate a bit, initially = 0
  - ❑ When page is referenced bit set to 1
  - ❑ Replace any with reference bit = 0 (if one exists)
    - ❑ We do not know the order, however
  - ❑ Basis for many LRU approximation algorithms

# LRU Approximation Algorithms


## ❑ Second-chance algorithm

- ❑ Generally FIFO, plus hardware-provided reference bit
- ❑ If page to be replaced has
  - ❑ Reference bit = 0 -> replace it
  - ❑ reference bit = 1 then:
    - ❑ set reference bit 0, leave page in memory
    - ❑ replace next page, subject to same rules

# Second-Chance (clock) Page-Replacement Algorithm



(a)



(b)

# Enhanced Second-Chance Algorithm

- ❑ Improve algorithm by using reference bit and modify bit (if available)
- ❑ Take ordered pair (reference, modify)
  - ❑ (0, 0) neither recently used nor modified – best page to replace
  - ❑ (0, 1) not recently used but modified – not quite as good, must write out before replacement
  - ❑ (1, 0) recently used but clean – probably will be used again soon
  - ❑ (1, 1) recently used and modified – probably will be used again soon and need to write out before replacement
- ❑ When page replacement called for, use the clock scheme but use the four classes replace page in lowest non-empty class
  - ❑ Might need to search circular queue several times

# Frame Allocation

- ❑ Each process needs minimum number of frames
  - ❑ E.g. IBM 370 – 6 pages to handle MVC instruction
- ❑ Maximum that a process can get of course is total frames in the system
- ❑ Two major allocation schemes
  - ❑ Fixed allocation
  - ❑ Priority allocation
- ❑ Many variations

# Fixed Allocation

- Equal allocation

- E.g. if there are 100 frames (after allocating frames for the OS) and 5 processes, give each process 20 frames

- Proportional allocation – allocate according to the size of process

- Dynamic as degree of multiprogramming, process sizes change

- $s_i$  = size of process  $p_i$

- $S = \sum s_i$

- $m$  = total number of frames

- $a_i$  = allocation for  $p_i = \frac{s_i}{S} \times m$



# Priority Allocation

- ❑ Use a proportional allocation scheme using priorities rather than size
- ❑ If process  $P_i$  generates a page fault,
  - ❑ select for replacement one of its frames
  - ❑ select for replacement a frame from a process with lower priority number


# Global vs. Local Allocation

- ❑ **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
  - ❑ A process cannot control its own page-fault rate
  - ❑ Greater throughput so more common
- ❑ **Local replacement** – each process selects from only its own set of allocated frames
  - ❑ More consistent per-process performance
  - ❑ But possibly underutilized memory

# Thrashing

- ❑ If a process does not have “enough” pages, the page-fault rate is very high
  - ❑ Page fault to get page
  - ❑ Replace existing frame
  - ❑ But quickly need replaced frame back
  - ❑ This leads to:
    - ❑ Low CPU utilization
    - ❑ Operating system thinking that it needs to increase the degree of multiprogramming
    - ❑ Another process added to the system
- ❑ **Thrashing**  $\equiv$  a process is busy swapping pages in and out

# Thrashing (Cont.)



# Demand Paging and Thrashing

- ❑ Why does demand paging work?

## Locality model

- ❑ Process migrates from one locality to another
  - ❑ Localities may overlap
- ❑ Why does thrashing occur?
  - $\Sigma$  size of locality > total memory size
  - ❑ Limit effects by using local or priority page replacement

# Working Set

## ❑ Main idea

- ❑ Define a working set as the set of pages in the most recent  $K$  page references to approximate the program's locality
- ❑ Keep the working set in memory will reduce page faults significantly

## ❑ Approximate working set

- ❑ The set of pages of a process used in the last  $T$  seconds


page reference table

... 2 6 1 5 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



# WSClock

- ☐ Follow the clock hand
- ☐ If the reference bit is 1
  - ☐ Set reference bit to 0
  - ☐ Set the current time for the page
  - ☐ Advance the clock hand
- ☐ If the reference bit is 0, check “time of last use”
  - ☐ If the page has been used within  $\delta$ , go to the next
  - ☐ If the page has not been used within  $\delta$  and modify bit is 1
    - ☐ Schedule the page for page out and go to the next
  - ☐ If the page has not been used within  $\delta$  and modify bit is 0
    - ☐ Replace this page



"WSCLOCK—a simple and effective algorithm for virtual memory management" by Richard W. Carr and John L. Hennessy, SOSP'81.